

R Implementation - Introduction to Double Robust Estimation for Causal Inference

Laura B. Balzer, PhD
University of Massachusetts-Amherst
lbalzer@umass.edu
Twitter: @LauraBBalzer

August 26, 2018

During the workshop, we will complete this R exercise together. After the workshop, an answer key will be distributed.

1 Preliminaries

1.1 Workshop Goals:

- (a.) Introduce the Roadmap for causal inference with an applied example (Petersen and van der Laan, 2014; Petersen and Balzer, 2014; Balzer et al., 2016).
- (b.) In R, implement the simple substitution estimator (a.k.a. parametric G-computation) and targeted maximum likelihood estimation (TMLE) with the ensemble algorithm Super Learner for estimation of the effect of a binary exposure.
- (c.) Show the `ltmle` and `drtmle` packages in R (Lendle et al., 2017; Benkeser, 2018).
- (d.) Time-permitting: Explore the performance of the estimators with simulations.

1.2 The Observed Data:

Load the data set `CausalWorkshop.csv` and assign it to object `data`.

```
> data<- read.csv("CausalWorkshop.csv")  
> # Hint 1. Do not type or copy the >. This represents the prompt in an R console.  
> # Hint 2. Press enter or return to evaluate the code.  
> # Hint 3. Hashtags are used in R for making comments. #YayTMLE!
```

Use the `tail` function to view the last 6 observations, and the `dim` function to obtain the dimensions of the data set:

```
> tail(data)
```

	W1	W2	W3	W4	A	Y
95	0	0.43866653	0.62099901	0.6736281	0	0.041108280

```

96  0 0.47763995  0.40945733  0.8789507 0 0.087362023
97  1 0.38601726  0.10916835 -0.9633303 0 0.045684975
98  1 0.51518289  1.02451436  1.6171894 1 0.000042700
99  0 0.05015548 -1.62360003 -1.1573370 0 0.005801751
100 1 0.88454275  0.05359586 -0.6562194 0 0.103371039

```

```
> dim(data)
```

```
[1] 100   6
```

The observed data consist of the following variables

- $W1$: Binary baseline covariate (values 0 and 1)
- $W2$: Continuous baseline covariate (range from 0 to 1)
- $W3$: Continuous baseline covariate (centered at 0)
- $W4$: Continuous baseline covariate (centered at 0)
- A : The exposure (1 for exposed; 0 for unexposed)
- Y : The outcome (range from 0 to 1).

Let $W = (W1, W2, W3, W4)$ represent the vector of baseline covariates (i.e. measured confounders). Then our observed data are

$$O = (W, A, Y)$$

with A as the exposure and Y as the outcome. There are $n = 100$ observations in the data set. We will use the following notation:

- \mathbb{P} : the distribution of the observed data $O = (W, A, Y)$
- $\mathbb{E}(Y|A, W)$: the conditional mean of the outcome, given the exposure and covariates
- $\mathbb{P}(A = 1|W)$: the conditional probability of being exposed, given the covariates (i.e. the propensity score)
- $\mathbb{P}(W)$: the marginal distribution of baseline covariates

In this workshop, we will implement three algorithms for estimating the G-Computation identifiability result (Robins, 1986):

$$\begin{aligned}\Psi(\mathbb{P}) &= \sum_w [\mathbb{E}(Y|A = 1, W = w) - \mathbb{E}(Y|A = 0, W = w)] \mathbb{P}(W = w) \\ &= \mathbb{E}[\mathbb{E}(Y|A = 1, W) - \mathbb{E}(Y|A = 0, W)]\end{aligned}$$

This is our statistical parameter, which is also called the estimand.

Before doing any estimation, it is useful to “set the seed” in R. This ensures that the same values are generated each time each time we run **all** of our code.

```

> set.seed(1)
> #More information about this function can be accessed with
> ?set.seed

```

2 The simple substitution estimator (i.e. parametric G-computation)

- **Step 1. Estimate the conditional mean outcome $\mathbb{E}(Y|A, W)$ with parametric regression.**

First, we run a regression of the outcome Y on the exposure A and covariates W . This provides an estimate the conditional mean outcome $\hat{\mathbb{E}}(Y|A, W)$. Here, we are using the “hat” notation to denote an estimate based on our observed data.

Suppose we were willing to assume the conditional expectation of the outcome Y is accurately described by the following function of the exposure A and covariates W :

$$\text{logit}[\mathbb{E}(Y|A, W)] = \beta_0 + \beta_1 A + \beta_2 W_1 + \beta_3 W_2 + \beta_4 W_3 + \beta_5 W_4$$

where $\text{logit}(x) = \log(x/1-x)$ is the log-odds.

- (a.) Use the `glm` function to fit the conditional mean function $\mathbb{E}(Y|A, W)$ with main terms parametric logistic regression.

```
> outcome.regression<- glm(Y~A+W1+W2+W3+W4, family='binomial', data=data)
```

The argument `family='binomial'` specifies logistic regression, and the argument `data=data` specifies the data set for fitting the regression.

Logistic regression is appropriate for a binary or bounded outcome in $[0,1]$ (e.g. a proportion or a probability). Using logistic regression over linear regression ensures that the bounds on the outcome are respected during the estimation process and provides robustness under strong confounding and/or rare outcomes. Therefore, ignore the following warning message.

Warning message:

```
In eval(expr, envir, enclos) : non-integer #successes in a binomial glm!
```

- (b.) We can see a summary of the output of regression fit with

```
> summary(outcome.regression)
>
```

• Step 2: Using the fitted regression, obtain the predicted outcomes.

Next, we obtain the expected (predicted) outcome for each observation under the exposure and given the covariates $\hat{\mathbb{E}}(Y|A=1, W)$ as well as the expected (predicted) outcome for each observation under no exposure and given the covariates $\hat{\mathbb{E}}(Y|A=0, W)$.

- (a.) Copy the data set `data` into two new data frames `data.txt` and `data.untxt`. Then set all units in `data.txt` to be treated/exposed ($A=1$) and all units in `data.untxt` to be untreated/unexposed ($A=0$). The columns of a data frame can be accessed with the `$` operator.

```
> # copying the observed data
> data.txt<- data.untxt <- data
> # set A=1 in the data.txt and A=0 in data.untxt
> data.txt$A <-1
> data.untxt$A <- 0

> # Verify with the colMeans function (which takes the mean in each data column)
> colMeans(data.txt)
> colMeans(data.untxt)
>
```

- (b.) Use the `predict` function to get the expected outcome for each observation i given the exposure and its covariates $\hat{\mathbb{E}}(Y_i|A_i=1, W_i)$.

```
> predict.outcome.txt<- predict(outcome.regression, newdata=data.txt, type='response')
```

The argument `newdata=data.txt` specifies that we want to predict the outcomes using as input the data set `data.txt`, where $A=1$ for all units. The argument `type='response'` specifies that we want the predictions returned on the original scale of the response (not the log-odds).

We have created a vector `predict.outcome.txt`, which contains the predicted outcomes for all units, given the exposure and their measured covariates. View the predicted outcomes for the first 6 units under the exposure with `head` function.

```
> head(predict.outcome.txt)
>
```

- (c.) Use the `predict` function to get the expected outcome for each observation i given no exposure and its covariates $\hat{\mathbb{E}}(Y_i|A_i = 0, W_i)$.

```
> predict.outcome.untxt<- predict(outcome.regression, newdata=data.untxt, type='response')
```

The argument `newdata=data.untxt` specifies that we want to predict the outcomes using as input `data.untxt`, where $A = 0$ for all units. As before, we want the predicted outcomes (i.e. `type='response'`).

```
> # Viewing the predicted outcomes for the first 6 units under no exposure:
> head(predict.outcome.untxt)
>
```

- **Step 3. Estimate the statistical parameter by substituting the predicted outcomes into the G-Computation formula.**

Finally, we estimate $\Psi(\mathbb{P})$ by averaging the difference in the predicted outcomes. This step standardizes (marginalizes) our estimates to the empirical distribution of confounders:

$$\begin{aligned}\Psi_{SimpleSubs}(\hat{\mathbb{P}}) &= \frac{1}{n} \sum_{i=1}^n \left[\hat{\mathbb{E}}(Y_i|A_i = 1, W_i) - \hat{\mathbb{E}}(Y_i|A_i = 0, W_i) \right] \\ &= \underbrace{\frac{1}{n} \sum_{i=1}^n \hat{\mathbb{E}}(Y_i|A_i = 1, W_i)}_{\text{Est. of } \mathbb{E}[\mathbb{E}(Y|A=1, W)]} - \underbrace{\frac{1}{n} \sum_{i=1}^n \hat{\mathbb{E}}(Y_i|A_i = 0, W_i)}_{\text{Est. of } \mathbb{E}[\mathbb{E}(Y|A=0, W)]}\end{aligned}$$

where $\hat{\mathbb{P}}$ is the empirical distribution, which places weight $1/n$ on each observation $O_i = (W_i, A_i, Y_i)$. The sample average is the non-parametric maximum likelihood estimator (NPMLE) of the covariate distribution.

```
> # estimate of E[E(Y|A=1,W)]
> mean(predict.outcome.txt)
> # estimate of E[E(Y|A=0,W)]
> mean(predict.outcome.untxt)
> # our point estimate of Psi(P)
> Simple.Subs <- mean(predict.outcome.txt - predict.outcome.untxt)
> Simple.Subs
```

- Consistency of the simple (non-targeted) substitution estimator depends on consistent estimation of outcome regression: $\mathbb{E}(Y|A, W)$. If our parametric regression model is incorrect, our point estimates can be biased and inference misleading.
- An estimator is *consistent* if the point estimates converge (in probability) to the estimand as sample size $n \rightarrow \infty$.

3 Targeted Maximum Likelihood Estimation (TMLE)

- **Step 1. Estimate the conditional mean outcome $\mathbb{E}(Y|A, W)$ with Super Learner.**

First, we use Super Learner to flexibly estimate the conditional expectation of the outcome Y , given the exposure A and covariates W .

- (a.) Use the `library` function to load the `SuperLearner` package (Polley et al., 2018).

```
> library("SuperLearner")
>
```

- (b.) Specify the Super Learner library with the following algorithms: main terms logistic regression (`SL.glm`), main terms logistic regression with all possible pairwise interactions (`SL.glm.interaction`), and generalized additive models (`SL.gam`) (Hastie and Tibshirani, 1990; Hastie, 2016). We are using this simple library for pedagogic purposes. In practice, we would want to use a larger library with a mixture of simple (e.g. parametric) and more aggressive libraries.

```
> SL.library<- c("SL.glm", "SL.glm.interaction", "SL.gam")
> # note these have to be in quotations
>
```

- (c.) As input to the `SuperLearner` algorithm, we need to specify the outcome-for-prediction, denoted Y , as well as the predictors, denoted X . In this case, we are interested in predicting the outcome Y with the measured covariates W and the exposure A . Therefore, in the following, we denote the outcome-for-prediction with `Y=data$Y` and the predictors as the baseline covariates and the exposure with `X=subset(data, select=-Y)`. The `subset` function is excluding the outcome Y from `data` and retaining (selecting) the covariates W and exposure A .

```
> SL.outcome.regression<- SuperLearner(Y=data$Y, X=subset(data, select=-Y),
                                     SL.library=SL.library, family="binomial")
```

We are also specifying the library and the appropriate family. The resulting object is called `SL.outcome.regression`. As before, ignore the warning message “In eval(expr, envir, enclos) : non-integer #successes in a binomial glm!”

```
> # We can examine the output by typing
> SL.outcome.regression
>
```

The `Risk` column gives the cross-validated estimates of the risk (i.e. expected loss) for each algorithm averaged across the 10 folds. The `Coef` column gives the weight of each algorithm in the convex combination.

• Step 2. Using the Super Learner fit, obtain the predicted outcomes.

Next, we use the Super Learner fit to obtain initial estimates of the expected (predicted) outcome for each observation given the observed exposure and the covariates $\hat{\mathbb{E}}(Y|A, W)$, given the treatment/exposure and covariates $\hat{\mathbb{E}}(Y|A = 1, W)$, and given no treatment/exposure and the covariates $\hat{\mathbb{E}}(Y|A = 0, W)$.

- (a.) Use the `predict` function to obtain initial estimates of the expected outcome, given the observed exposure and covariates $\hat{\mathbb{E}}(Y|A, W)$.

```
> SL.predict.outcome<- predict(SL.outcome.regression, newdata=data)$pred
```

The argument `newdata=data` specifies that we want to predict the outcome using as input the observed exposure and covariates. We have created a vector `SL.predict.outcome`, which contains initial estimates of the expected outcome, given the observed exposure and covariates $\hat{\mathbb{E}}(Y|A, W)$.

```
> head(SL.predict.outcome)
>
```

- (b.) Also obtain the initial estimates of the expected outcome for all units under the treatment/exposure $\hat{\mathbb{E}}(Y|A = 1, W)$. Now we specify `newdata=data.txt` to predict the outcome using as input `data.txt`, where $A = 1$ for all units.

```
> SL.predict.outcome.txt<- predict(SL.outcome.regression, newdata=data.txt)$pred
> head(SL.predict.outcome.txt)
>
```

- (c.) Finally, obtain the initial estimates of the expected outcome for all units under no treatment/exposure $\hat{\mathbb{E}}(Y|A = 0, W)$. Now we specify `newdata=data.untxt` to predict the outcome using as input `data.untxt`, where $A = 0$ for all units.

```
> SL.predict.outcome.untxt<- predict(SL.outcome.regression, newdata=data.untxt)$pred
> head(SL.predict.outcome.untxt)
>
```

• **Step 3: Estimate the propensity score $\mathbb{P}(A = 1|W)$ with Super Learner.**

Now we want to use Super Learner to flexibly estimate the conditional probability of being exposed, given the measured covariates.

- (a.) As input to the `SuperLearner` algorithm, we need to specify the outcome-for-prediction as well as the predictors. In this case, we are interested in predicting the exposure A with the measured covariates W . Therefore, in the following, we denote the “outcome” with `Y=data$A` and the predictors as the baseline covariates with `X=subset(data, select=-c(A,Y))`. Here, the `subset` function is excluding both the exposure and outcome (A, Y) from `data` and retaining (selecting) the covariates W .

```
> SL.pscore<- SuperLearner(Y=data$A, X=subset(data, select=-c(A,Y)),
                           SL.library=SL.library, family="binomial")
```

For simplicity, we are using the same library, although this is not a requirement. We are telling `SuperLearner` that the “outcome” (here, the exposure) is binary (`family="binomial"`). We are calling the resulting object `SL.pscore`.

```
> # We can examine the output by typing
> SL.pscore
>
```

• **Step 4: Using the Super Learner fit, create the clever covariate $\hat{H}(A, W)$:**

$$\hat{H}(A, W) = \frac{\mathbb{I}(A = 1)}{\hat{\mathbb{P}}(A = 1|W)} - \frac{\mathbb{I}(A = 0)}{\hat{\mathbb{P}}(A = 0|W)}.$$

The clever covariate for exposed units is 1 over the predicted probability of being treated/exposed, and the clever covariate for unexposed units is -1 over the predicted probability of not being treated/exposed.

- (a.) We can access the estimated propensity score $\hat{\mathbb{P}}(A_i = 1|W_i)$ with

```
> SL.predict.prob.txt<- SL.pscore$SL.predict
> # Note this is equivalent to
> # predict(SL.pscore, newdata=data)$pred
> summary(SL.predict.prob.txt)
> # In practice, we may need to bound the predicted propensity scores away from (0,1).
>
```

- (b.) Likewise, we can calculate the predicted probability of not being treated/exposed, given the baseline covariates.

```
> SL.predict.prob.untxt<- 1- SL.predict.prob.txt
```

- (c.) Use these estimates to create the clever covariate:

$$\hat{H}(A, W) = \left(\frac{\mathbb{I}(A = 1)}{\hat{\mathbb{P}}(A = 1|W)} - \frac{\mathbb{I}(A = 0)}{\hat{\mathbb{P}}(A = 0|W)} \right).$$

```
> H.AW<- as.numeric(data$A==1)/SL.predict.prob.txt -
          as.numeric(data$A==0)/SL.predict.prob.untxt
> summary(H.AW)
>
```

- (d.) Also evaluate the clever covariate at $A = 1$ and $A = 0$ for all units:

$$\hat{H}(1, W) = \frac{1}{\hat{\mathbb{P}}(A = 1|W)} \quad \text{and} \quad \hat{H}(0, W) = \frac{-1}{\hat{\mathbb{P}}(A = 0|W)}$$

Call the resulting values `H.1W` and `H.0W`, respectively.

```
> H.1W<- 1/SL.predict.prob.txt
> H.0W<- -1/SL.predict.prob.untxt
```

```
> # We have created 3 new vectors H.AW, H.1W, H.0W.
> # viewing the last six observations in each vector along with the exposure status:
> tail(data.frame(data$A, H.AW, H.1W, H.0W))
>
```

- **Step 5. Target the initial estimator of the conditional mean outcome $\hat{\mathbb{E}}(Y|A, W)$ with information in the estimated propensity score $\hat{\mathbb{P}}(A = 1|W)$.**

- (a.) Run a univariate regression of the outcome Y on the clever covariate $\hat{H}(A, W)$ with the initial estimator as offset. Specifically for a binary or bounded continuous outcome, we estimate the coefficient ϵ by fitting the following logistic regression model

$$\text{logit}[\hat{\mathbb{E}}^*(Y|A, W)] = \text{logit}[\hat{\mathbb{E}}(Y|A, W)] + \epsilon \hat{H}(A, W).$$

Note there is no intercept (i.e. there is no β_0 term), and the coefficient on the (*logit*) of the initial estimator is set to 1.

```
> logitUpdate<- glm(data$Y ~ -1 +offset(qlogis(SL.predict.outcome)) + H.AW,
  family='binomial')
```

- We are again calling the `glm` function to fit a generalized linear model.
- On the left hand side of the formula, we have the outcome Y .
- On the right hand side of the formula, we suppress the intercept by including `-1` and use as `offset` the *logit* of our initial Super Learner estimates `SL.predict.outcome`. In R, *logit*(x) = $\log(x/(1-x))$ function is given by `qlogis(x)`.
- The only main term in the regression is the clever covariate $\hat{H}(A, W)$.
- Including `family='binomial'` runs logistic regression.
- Again ignore any warning message.

```
> # we can examine the output by typing
> summary(logitUpdate)
>
```

- (b.) Let `epsilon` denote the resulting maximum likelihood estimate of the coefficient on the clever covariate `H.AW`.

```
> epsilon<- logitUpdate$coef
> epsilon
```

- (c.) Plug-in the estimated coefficient $\hat{\epsilon}$ to yield **targeted** estimates of the expected outcome under the exposure $\hat{\mathbb{E}}^*(Y|A = 1, W)$ and under no exposure $\hat{\mathbb{E}}^*(Y|A = 0, W)$:

$$\begin{aligned}\hat{\mathbb{E}}^*(Y|A = 1, W) &= \text{logit}^{-1} \left[\text{logit}[\hat{\mathbb{E}}(Y|A = 1, W)] + \hat{\epsilon} \hat{H}(1, W) \right] \\ \hat{\mathbb{E}}^*(Y|A = 0, W) &= \text{logit}^{-1} \left[\text{logit}[\hat{\mathbb{E}}(Y|A = 0, W)] + \hat{\epsilon} \hat{H}(0, W) \right]\end{aligned}$$

where logit^{-1} is the inverse-*logit*, $\hat{H}(1, W)$ is the clever covariate evaluated for all units under the exposure, and $\hat{H}(0, W)$ is the clever covariate evaluated for all units under no exposure.

```
> target.predict.outcome.txt<- plogis( qlogis(SL.predict.outcome.txt)+ epsilon*H.1W)
> target.predict.outcome.untxt<- plogis( qlogis(SL.predict.outcome.untxt)+ epsilon*H.0W)
>
```

In R, the inverse-*logit* function is given by `plogis(x)`.

- **Step 6. Estimate the statistical parameter by substituting the targeted predictions into the G-Computation formula.**

Estimate $\Psi(\mathbb{P})$ by averaging the difference in the targeted predictions:

$$\Psi_{TMLE}(\hat{\mathbb{P}}) = \frac{1}{n} \sum_{i=1}^n \left[\hat{\mathbb{E}}^*(Y_i|A_i = 1, W_i) - \hat{\mathbb{E}}^*(Y_i|A_i = 0, W_i) \right]$$

```
> TMLE<- mean(target.predict.outcome.txt- target.predict.outcome.untxt)
> TMLE

> # we can also estimate  $E[E(Y|A=1,W)]$ 
> mean(target.predict.outcome.txt)
> # and estimate  $E[E(Y|A=0,W)]$ 
> mean(target.predict.outcome.untxt)
```

- Compare the point estimates from the 3 methods:

```
> c(Simple.Subs, TMLE)
```

The true value is -2.14%.

- TMLE is double robust; it will be consistent if *either* the conditional mean outcome $\mathbb{E}(Y|A, W)$ or the propensity score $\mathbb{P}(A = 1|W)$ is estimated consistently.
- If both $\mathbb{E}(Y|A, W)$ and $\mathbb{P}(A = 1|W)$ are estimated consistently (at a fast enough rate), the TMLE will be efficient and achieve the lowest possible asymptotic variance over a large class of (semi-parametric) estimators.
- These asymptotic properties describe what happens when sample size goes to infinity and also typically translate into lower bias and variance in finite samples.
- *Note:* If you call `SuperLearner` multiple times, you might get slightly different output. This is due, in part, to working with a limited library of algorithms in which no one algorithm dominates in performance. Furthermore, during the cross-validation step, the observations are randomly allocated to folds and with a relatively small sample size ($n = 100$), this random allocation could impact the weight distribution (`Coef`) for the algorithms.

4 The basics of the ltmle package

The `ltmle` package (Lendle et al., 2017) expands the previous `tmle` package (Gruber and van der Laan, 2012). Specifically `ltmle` estimates parameters corresponding to point-treatment exposures, longitudinal exposures, marginal structural working models, dynamic treatment regimes and much more!

- **Step 1. Load the SuperLearner and ltmle packages.**

```
> library('SuperLearner')
> library('ltmle')
> # we can learn a lot more about the function by reading the help file
> ?ltmle
```

- The basic input to the function is the data set `data`, the exposure variable(s) `Anodes`, the outcome(s) `Ynodes`, and the exposure levels of interest `abar`.
 - The user can also specify censoring variables `Cnodes`, time-dependent covariates `Lnodes`, weights `observation.weights`, and the independent unit `id`. (See the help file for more information.)
 - Initial estimates of the conditional mean outcome $\mathbb{E}(Y|A, W)$ can be estimated according to a user-specified regression formula `Qform` or estimated with Super Learner `SL.library`.
 - Initial estimates of the propensity score $\mathbb{P}(A = 1|W)$ can be estimated according to a user-specified regression formula `gform` or estimated with Super Learner `SL.library`.
- **Step 2. Call the ltmle function using Super Learner to estimate the conditional mean outcome $\mathbb{E}(Y|A, W)$ and the propensity score $\mathbb{P}(A = 1|W)$.**


```
> ltmle.package<- ltmle(data=data, Anodes='A', Ynodes='Y', abar=list(1,0),
  SL.library=SL.library)
```

Here, `abar=list(1,0)` specifies the comparison of interest: all exposed ($A = 1$) vs. all unexposed ($A = 0$).

- **Step 3. Use the summary function to obtain point estimates and get inference.**

```
> summary(ltmle.package)
```

5 Bonus: drtmle package

- The `drtmle` package is an alternative and can handle categorical exposures (and provides valid statistical inference for IPW with Super Learner) (Benkeser, 2018).
- For more information and an example with multi-level exposures, check out: https://benkeser.github.io/drtmle/articles/using_drtmle.html#multi-level-treatments

References

- L. Balzer, M. Petersen, and M.J. van der Laan. Tutorial for causal inference. In P. Buhlmann, P. Drineas, M. Kane, and M. van der Laan, editors, *Handbook of Big Data*. Chapman & Hall/CRC, 2016.
- David Benkeser. *drtmle: Doubly-Robust Nonparametric Estimation and Inference*, 2018. URL <https://CRAN.R-project.org/package=drtmle>. R package version 1.0.3.
- S. Gruber and M.J. van der Laan. tmle: An R package for targeted maximum likelihood estimation. *Journal of Statistical Software*, 51(13):1–35, 2012. doi: 10.18637/jss.v051.i13.
- T. Hastie. *gam: Generalized Additive Models*, 2016. URL <http://CRAN.R-project.org/package=gam>. R package version 1.14.
- T.J. Hastie and R.J. Tibshirani. *Generalized additive models*. Chapman & Hall, Boca Raton, 1990.
- S.D. Lendle, J. Schwab, M.L. Petersen, and M.J. van der Laan. ltmle: An R package implementing targeted minimum loss-based estimation for longitudinal data. *Journal of Statistical Software*, 81(1):1–21, 2017.
- M.L. Petersen and L.B. Balzer. Introduction to Causal Inference, 2014. URL www.ucbbiostat.com.
- M.L. Petersen and M.J. van der Laan. Causal models and learning from data: Integrating causal modeling and statistical estimation. *Epidemiology*, 25(3):418–426, 2014.
- M.L. Petersen, K.E. Porter, S. Gruber, Y. Wang, and M.J. van der Laan. Diagnosing and responding to violations in the positivity assumption. *Statistical Methods in Medical Research*, 21(1):31–54, 2012. doi: 10.1177/0962280210386207.
- E. Polley, E. LeDell, C. Kennedy, and M. van der Laan. *SuperLearner: Super Learner Prediction*, 2018. URL <http://CRAN.R-project.org/package=SuperLearner>. R package version 2.0-24.
- J.M. Robins. A new approach to causal inference in mortality studies with sustained exposure periods—application to control of the healthy worker survivor effect. *Mathematical Modelling*, 7:1393–1512, 1986. doi: 10.1016/0270-0255(86)90088-6.